

# AMIT RAJ REDDY DHARAM

+1 903-990-6812 | reddy.amitraj@gmail.com | linkedin.com/in/arrd19/ | github.com/Ames-Zero | Tempe, AZ

## SUMMARY

Computer Science master's student with 2 years of hands-on machine learning engineering expertise, focusing on deep learning model development using PyTorch, data preparation pipelines, and production ML systems. Proven expertise in Python, NLP, computer vision, GenAI and cloud-based AI systems.

## EDUCATION

<b>Masters in Computer science</b> , Arizona State University	May 2026
Relevant Coursework: Artificial Intelligence, Statistical Machine Learning, Data Mining, Cloud Computing, NLP	Tempe, AZ
<b>Bachelors in Computer Science and MSc in Mathematics</b> , BITS Pilani	May 2023
Relevant Coursework: Data Structures and Algorithms (C/C++), OOP (Java), Deep Learning (Python), DBMS	India

## TECHNICAL SKILLS

**Languages and Technologies:** Python, Java, JavaScript, C/C++, Flask, FastAPI, MySQL, MongoDB, Redis, Neo4j  
**ML frameworks and libraries:** PyTorch, TensorFlow, Pandas, PySpark, Scikit-Learn, NumPy, Matplotlib, Open3D, XGboost  
**Natural Language Processing (NLP):** Hugging Face Transformers, LangChain, FAISS, Pinecone, BERT, OpenAI APIs  
**Developer tools:** Amazon Web Services (AWS), Microsoft Azure, Docker, Kubernetes, CircleCI, Github, Terraform  
**Concepts:** Deep learning, AI cloud deployment, backend development and API design, RAG, GenAI, computer vision

## PROFESSIONAL EXPERIENCE

<b>AI Data Science Analyst(Part-time)</b>   Arizona State University   Tempe, AZ	Sept 2024 – June 2025
• Implemented the backend for a NL2SQL conversational agent using Python, FastAPI allowing users to query, filter and display data from multi-table databases using conversational inputs, with query execution accuracy over 85%	
• Optimized natural language to SQL module responses by incorporating team lead and engineer feedback into prompt engineering, fine-tuned LLMs, and GraphRAG implementation, leading in a 5% increase in accuracy	
• Integrated LangChain for fast orchestration and Neo4j for efficient semantic search using embeddings; utilized AWS EKS for scalable inference; and reduced response time by 30% over 100+ questions per session.	
<b>Software Development Engineer- ML</b>   TELUS Digital   India	July 2023 – July 2024
• Instruction-tuned open-source LLMs with PyTorch on scientific topics using PEFT-based methods for synthetic data creation using a RAG pipeline. Increased manual question generation throughput by 50%, reducing generation time by 40%	
• Mentored an intern to test the LLM against conventional NLP benchmarks (MATH, MMLU) and internal domain-specific datasets, achieving a +6% improvement in benchmark scores and 50% faster evaluation timeframes	
• Evaluated a custom trained 3D object detection (Voxel-based) and tracking models(DeepSORT) on the KITTI benchmark, achieving around 90% precision for object detection and 70% MOTA (Multi Object Tracking Accuracy) for tracking	
<b>Software Development Intern - ML</b>   TELUS Digital   India	Aug 2022 – July 2023
• Prototyped DeepSORT and ByteTrack algorithms for bounding box tracking over video sequences resulted in 80% mAP object recognition and 75% IDF1 multi-object tracking accuracy, as confirmed against MOT17 benchmarks	
• Improved image segmentation precision by 23% across 15,000+ annotated frames by uptraining U-Net and Mask R-CNN architectures using the CityScapes dataset, reducing manual annotation time from 8 hours to 5 hours per dataset batch	
• Deployed ML APIs serving 15+ active projects with sub-200 ms response times by establishing FastAPI and TorchServe endpoints, containerizing the segmentation models with Docker, and hosting on AWS EC2 with JWT authentication	

## PROJECTS

<b>WardrobeAI (Reccomendations and Image Generation)</b>   Winner Sunhacks, Sept 2025	Project-Link
• Built a Python backend that integrates with MongoDB and AWS S3, linking Gemini chat API and Gemini Nano-Banana models across 5+ microservices to provide virtual try-ons and outfit recommendations for over 3000 clothing items	
• Coordinated a 4-person team to architect and deploy a LinUCB based recommendation engine with conversational chatbot, generating personalized outfit suggestions for 10+ occasions like birthdays, weddings etc.	
<b>Tom-Riddle's Diary (AI Chatbot agent)</b>   Winner InnovationHacks, Apr 2025	Project-Link
• Engineered and integrated a python, FastAPI backend using ChromaDB to enable long-term semantic memory and efficient diary entry retrieval, reducing response time by 35% and supporting contextual AI reflections	
• Integrated Google Gemini LLMs into a conversational backend pipeline, enabling natural language queries and persona based tone generation handling over 1,000 personalized interactions with >90% response relevance	